

MCP: using Java and Quarkus to bridge LLMs with your applications and data

Horacio Gonzalez

2025-04-02



@LostInBrittany



clever cloud



SVQJUG

Summary




1. Introduction
2. Understanding LLMs & Their Landscape
3. Making Java applications LLM-aware
4. Model Context Protocol (MCP): The missing link
5. Live demo: Building an MCP server in Java with Quarkus
6. The future of Java & LLM integration
7. Q&A and discussion



Introduction

LLMs are changing software development—
how can Java developers take full advantage?




Horacio Gonzalez

@LostInBrittany

Spaniard Lost in Brittany


Head of DevRel



clever cloud



clever cloud





clever cloud

From Code to Product

Our mission: give more **speed** to your **teams**
and better **quality** to your **projects**




clever cloud

@LostInBrittany



Why are we talking about this?




LLMs are changing development,
but most Java apps don't fully leverage them

How do you use LLMs for your dev job?



1. Who here has already used LLM?
2. Who here has already used LLM professionally?
3. Who here has already used LLM with code?
4. Who here has already used LLMs in a Java app?



How LLMs are changing dev jobs



The 70% problem: Hard truths about AI-assisted coding

A field guide and why we need to rethink our expectations



ADDY OSMANI

DEC 04, 2024



800



44



121

Share

After spending the last few years embedded in AI-assisted development, I've noticed a fascinating pattern. While engineers report being dramatically more productive with AI, the actual software we use daily doesn't seem like it's getting noticeably better. What's going on here?


I think I know why, and the answer reveals some fundamental truths about software development that we need to reckon with. Let me share what I've learned.

<https://addyo.substack.com/p/the-70-problem-hard-truths-about>



Understanding LLMs & Their Landscape

Closed-source, open-source, local—
choosing the right LLM for your Java application.



LLMs come in different flavors



Not all LLMs are created equal

They have different trade-offs
in capabilities, accessibility, and control



Choosing the right one depends on
your use case, security needs, and infrastructure.

Closed-source LLMs (Cloud-based APIs)



Examples

- OpenAI (ChatGPT), Anthropic (Claude), Google (Gemini), Microsoft (Copilot)

Advantages:

- Powerful and well-trained (best models available)
- Easy to use via APIs
- Regularly updated & improved

Challenges:

- Black box (you don't control how they work)
- Expensive (API calls can add up quickly)
- Data privacy concerns (sending requests to external servers)

When to use?

- If you need the most advanced models and don't mind API costs or external dependencies.



clever cloud

Open-source LLMs (Self- or cloud-hosted)



Examples

- Meta's Llama 3, Mistral, Google's Gemma, Alibaba's Qwen



Advantages:

- Greater control (you know exactly how the model works)
- Can be fine-tuned for specific needs
- No external API costs



Challenges:

- Requires more setup (you have to run the model yourself)
- May not be as powerful as the latest closed models
- Needs infrastructure (e.g., GPUs for hosting)



When to use?

- If you need control over the model & lower costs but are okay with slightly weaker performance



Local models (on your machine or server)



Examples

- Ollama, GGUF-based models (e.g., Llama, Mistral, Mixtral)

Advantages:

- Works offline (great for security-sensitive applications)
- No API costs (completely free to use once set up)
- Low latency (responses are instant if hardware is good)

Challenges:

- Limited by your hardware (needs a strong CPU/GPU)
- Not always as capable as cloud-hosted models
- Setup complexity (installing and optimizing models)

When to use?

- If you need privacy and control, and you have the hardware to run an LLM efficiently



Llama 3



Gemma





clever cloud

Choosing the Right Model for your Apps



- Cloud APIs
 - Great for rapid development, but costly and not always secure
- Self-hosted open models
 - Best balance for long-term control and scalability
- Local models
 - Best for privacy-sensitive applications



clever cloud

@LostInBrittany





Making Java Applications LLM-Aware

LangChain4j simplifies LLM integration—
let's see how it works!



clever cloud

@LostInBrittany



Two ways to integrate LLMs with Java



1. Java applications using LLMs

- Using LLMs as assistants, API consumers, or reasoning engines
- Easier and works well for code generation, chatbots, and AI assistants

2. LLMs using Java applications

- Exposing Java functions, APIs, and databases to LLMs for tool execution
- More powerful but requires tool calling and execution control.

What is LangChain4j?



LangChain4j is a Java library that simplifies LLM integration

It's the Java equivalent of LangChain (Python/JS)

- ✓ Prompt management – Structured input for LLMs.
- ✓ Memory – Keeping conversational state.
- ✓ Agents – Allowing LLMs to decide which tools to use.
- ✓ Tool calling – Exposing functions to LLMs.



LangChain4j Hello Sevilla JUG



```
HelloLangChain4j01.java


///usr/bin/env jbang "$0" "$@" ; exit $?
//DEPS dev.LangChain4j:LangChain4j-open-ai:1.0.0-beta1
import dev.LangChain4j.model.openai.OpenAiChatModel;
import static dev.LangChain4j.model.openai.OpenAiChatModelName.GPT_4_O_MINI;


public class LangChain4jHelloMadridJug {
    public static void main(String... args) {
        String apiKey = System.getenv("OPENAI_API_KEY");
        OpenAiChatModel model = OpenAiChatModel.builder()
            .apiKey(apiKey).modelName(GPT_4_O_MINI).build();
        String response = model.chat(
            "I am doing a demo at Sevilla JUG. Can you introduce yourself and say hello?");
        System.out.println(response);
    }
}
```


With Jbang, of course


We want to see it working!





 **langchain4j-for-tools-and-mcp-demo** Public


 Pin


 Unwatch 1


 Fork 0


 Star 0


 main


 1 Branch






 0 Tags


 Go to file


 Add file

 Code

 **LostInBrittany** Updating the README d28fa1f · 1 minute ago 6 Commits

 01-Hello-LangChain4j	Refactoring and adding the README	2 hours ago
 02-Lying-weather-Tool	Adding Real Weather Demo	9 minutes ago
 03-Real-weather-Tool	Adding Real Weather Demo	9 minutes ago
 LICENSE	Initial commit	3 hours ago
 README.md	Updating the README	1 minute ago

 README

 MIT license


LangChain4j for Tools and MCP - Demo Repository


This repository contains all the code and examples from my talk:
MCP: using Java and Quarkus to bridge LLMs with your applications and data, presented at:


- 2024-03-05 - [Madrid JUG](#) (Madrid, Spain) - [Slides](#)


About


A series of demos to illustrate my talk "MCP: using Java and Quarkus to bridge LLMs with your applications and data"


 Readme

 MIT license

 Activity

 0 stars

 1 watching

 0 forks

Releases

No releases published
[Create a new release](#)

Packages

No packages published
[Publish your first package](#)

Languages




<https://github.com/LostInBrittany/langchain4j-for-tools-and-mcp-demo/>



What are tools in LLMs?




Tools allow LLMs to do **more than just generate text**. They can interact with APIs, databases, and execute functions.



Instead of returning "I don't know," an LLM can call a tool **to fetch real data**.



An LLM without Tools can't answer this



What's the weather like in Sevilla today?

I'm unable to provide real-time
information or current weather updates.



Example of how tool calling works




Weather API



`getWeather("Sevilla (ES)")`

`{"weather": "sunny",
"temperature": "1.8°C"}`



What's the weather
like in Sevilla today?

Today it is sunny in
Sevilla, but very cold,
take a coat.




LLM recognizes it needs an external function and calls it.
It integrates the result into a natural-language response.

Why this matters?



- Moves LLMs from static text generation
 - dynamic system components
- Increases accuracy & real-world usability
- Allows developers to control what the LLM can access



What's the weather like
in Madrid today?

Today it is sunny in Madrid,
but very cold, take a coat.



clever cloud

@LostInBrittany



Using LangChain4j to Define LLM Tools



```

LyingWeatherTool.java

//DEPS dev.langchain4j:langchain4j:1.0.0-beta1

import dev.langchain4j.agent.tool.Tool;

public class LyingWeatherTool{
    @Tool("A tool to get the current weather in a city")
    public static String getWeather(String city) {
        return "The weather in " + city + " is sunny and hot.";
    }
}

```

LangChain4j gives as a Tool framework

Using LangChain4j to Define LLM Tools




```
LyingWeatherTool.java


[...]  
    Assistant assistant = AIServices.builder(Assistant.class)  
        .chatLanguageModel(model).chatMemory(chatMemory)  
        .tools(new LyingWeatherTool()).build();  
    System.out.println("-----");  
    String question = "What will the weather be like in Madrid tomorrow?";  
    String response = assistant.chat(question);  
    System.out.println(response);  
    System.out.println("-----");  
[...]
```


And an AI Service abstraction


Let's see the code!





**langchain4j-for-tools-and-mcp-demo** Public


 Pin


 Unwatch 1


 Fork 0


 Star 0


 main

 1 Branch

 0 Tags

 Add file

 Code

**LostInBrittany** Updating the README d28fa1f · 1 minute ago 6 Commits

01-Hello-LangChain4j	Refactoring and adding the README	2 hours ago
02-Lying-weather-Tool	Adding Real Weather Demo	9 minutes ago
03-Real-weather-Tool	Adding Real Weather Demo	9 minutes ago
LICENSE	Initial commit	3 hours ago
README.md	Updating the README	1 minute ago

README

MIT license


LangChain4j for Tools and MCP - Demo Repository


This repository contains all the code and examples from my talk:
MCP: using Java and Quarkus to bridge LLMs with your applications and data, presented at:


- 2024-03-05 - [Madrid JUG](#) (Madrid, Spain) - [Slides](#)


About


A series of demos to illustrate my talk "MCP: using Java and Quarkus to bridge LLMs with your applications and data"


 Readme

 MIT license

 Activity

 0 stars

 1 watching

 0 forks

Releases


No releases published
[Create a new release](#)

Packages


No packages published
[Publish your first package](#)

Languages

<https://github.com/LostInBrittany/langchain4j-for-tools-and-mcp-demo/>

**clever cloud**

@LostInBrittany



Model Context Protocol (MCP): The missing link



MCP bridges LLMs with your applications,
enabling controlled, real-world interactions



clever cloud



Why Do We Need MCP?



Function calling is powerful,
why do I need another concept?



LLM function calling is useful, but lacks structure

Why Do We Need MCP?



Problem

- LLMs **don't automatically know** what functions exist.
- **No standard way** to expose an application's capabilities.
- **Hard to control** security and execution flow.

Function calling is powerful,
why do I need another concept?



Solution: MCP

- MCP **defines a standard way** to describe and expose functionalities.
- **Applications stay in control** over what LLMs can do.



Model Context Protocol



Model Context Protocol

Search...

GitHub >

Documentation

- Python SDK
- TypeScript SDK
- Kotlin SDK
- Specification

Get Started

- Introduction
- Quickstart >
- Example Servers
- Example Clients

Tutorials

- Building MCP with LLMs
- Debugging
- Inspector

Get Started

Introduction

Get started with the Model Context Protocol (MCP)

Kotlin SDK released! Check out what else is new.

MCP is an open protocol that standardizes how applications provide context to LLMs. Think of MCP like a USB-C port for AI applications. Just as USB-C provides a standardized way to connect your devices to various peripherals and accessories, MCP provides a standardized way to connect AI models to different data sources and tools.

Why MCP?

MCP helps you build agents and complex workflows on top of LLMs. LLMs frequently need to integrate with data and tools, and MCP provides:

- A growing list of pre-built integrations that your LLM can directly plug into
- The flexibility to switch between LLM providers and vendors
- Best practices for securing your data within your infrastructure

On this page

Why MCP?

- General architecture

Get started

- Quick Starts
- Examples

Tutorials

- Explore MCP
- Contributing
- Support and Feedback




De facto standard for exposing system capabilities to LLMs

<https://modelcontextprotocol.io/>

How MCP works



- Applications define an MCP manifest (structured JSON).
- The manifest describes available functions, input/output formats, and security policies.
- LLMs can discover and request function execution safely.



Weather
MCP Server



MCP is provider-agnostic



Works with any LLM provider



Ensures standardized function exposure across platforms

Understanding the MCP Manifest



```
mcp-manifest.json

{ "functions": [
  {
    "name": "getWeather",
    "description": "Fetches the current weather for a city.",
    "parameters": {
      "city": { "type": "string", "required": true },
      "countryCode": { "type": "string", "required": true }
    }
  }
]
}
```

- Lists available tools
- Describes expected inputs/outputs
- Defines execution policies

That's all, folks!

Thank you all!



clever cloud

@LostInBrittany

