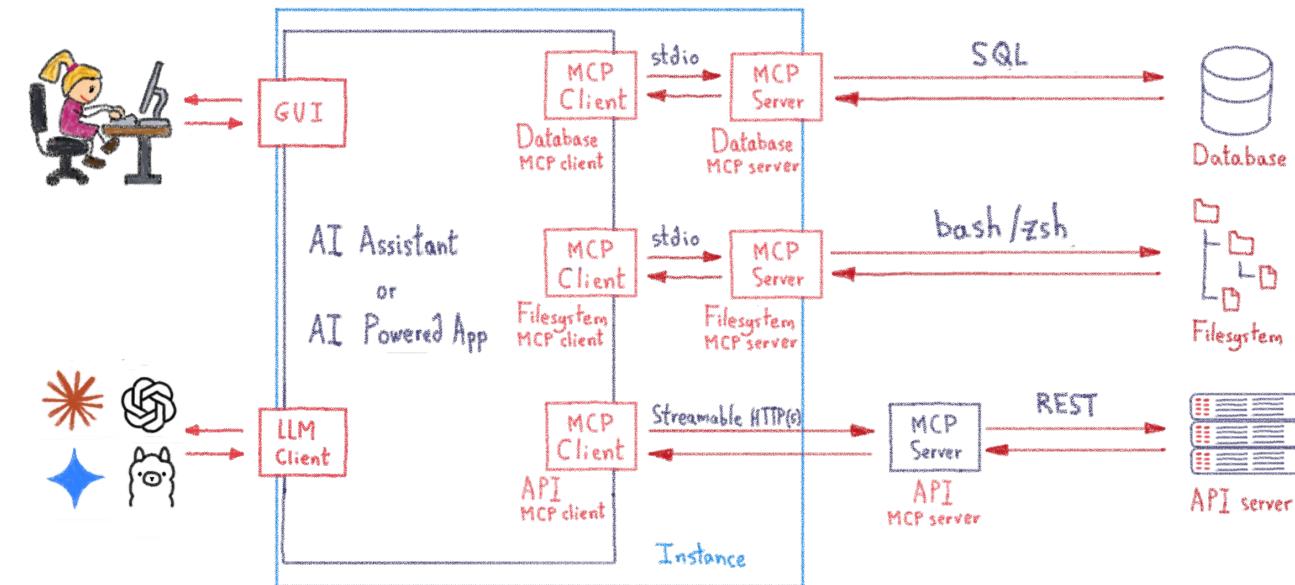
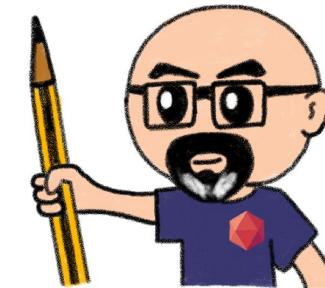


# Introducción al MCP: Conectando LLMs con tus Aplicaciones y Datos...

Horacio Gonzalez

2025-05-21

@LostInBrittany



# Horacio Gonzalez

@LostInBrittany

Spaniard Lost in Brittany

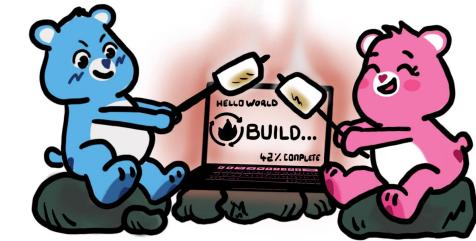
Head of DevRel



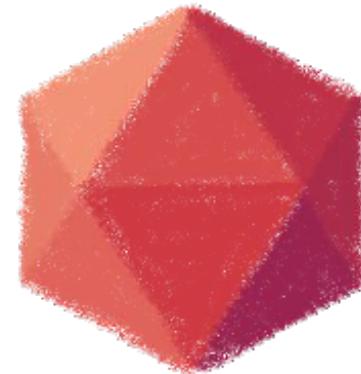
clever cloud



Finist  
Devs



@LostInBrittany



clever cloud

From Code to Product



Our mission: give more **speed** to your **teams**  
and better **quality** to your **projects**



# Summary

1. Introduction
2. LLM evolution
3. Model Context Protocol (MCP)
4. Architecture of MCP
5. MCPs are APIs
6. Q&A and discussion

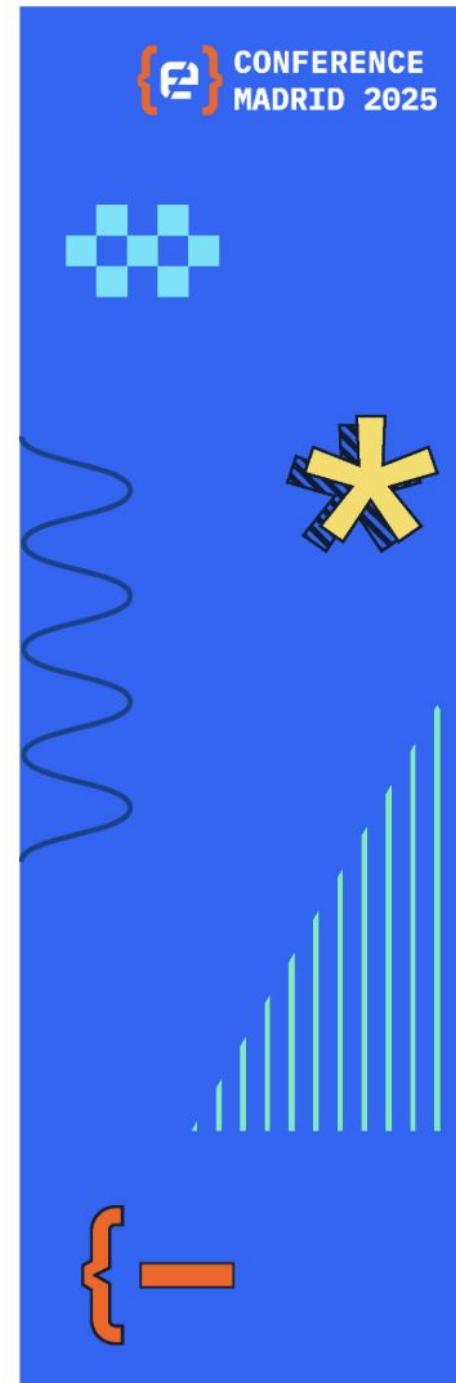


# Introduction

LLMs are changing software development, they say...  
how about you?



OR



# Why are we talking about this?



LLMs are changing development,  
but individual devs don't always leverage them

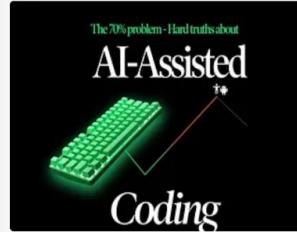
# How do you use LLMs for your dev job?

1. Who here has already used LLM?
2. Who here has already used LLM professionally?
3. Who here has already used LLM to assist with code?
4. Who here has already used LLMs by coding?



# How LLMs are changing dev jobs

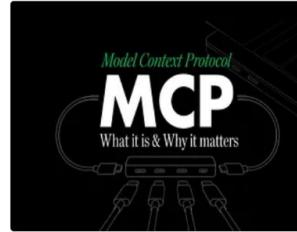
## Featured AI articles



The 70% Problem in AI Coding



Maximizing the Human 30% of AI-Assisted Coding



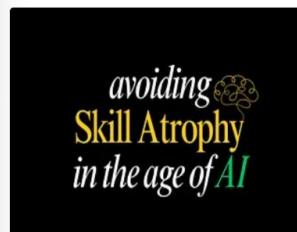
MCP: The Model Context Protocol



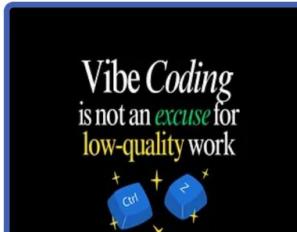
Why I Use Cline for AI Engineering



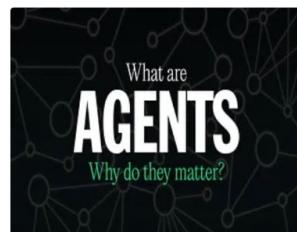
Leading Effective ENGINEERING TEAMS in the age of GenAI



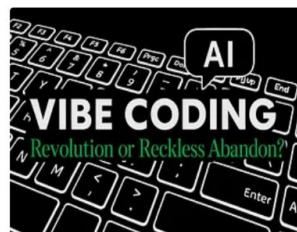
Avoiding Skill Atrophy in the Age of AI



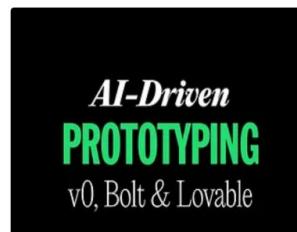
Vibe Coding is not an excuse for low-quality work



What are AI Agents? why do they matter?



Vibe Coding: Revolution or Reckless Abandon?



AI-Driven PROTOTYPING v0, Bolt, and Lovable Compared

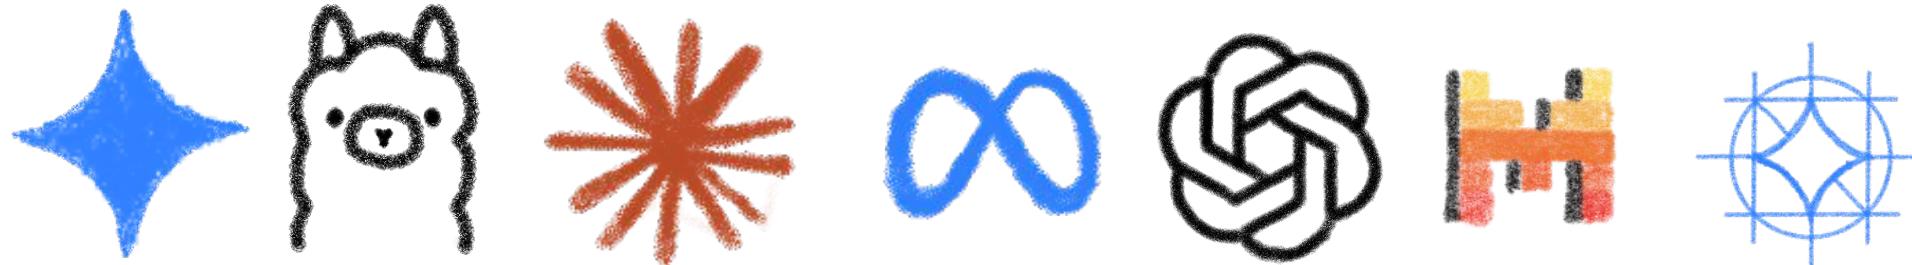
A point of view I find balanced: Addy Osmani

<https://addyosmani.com/>

# LLMs come in different flavors

Not all LLMs are created equal

They have different trade-offs  
in capabilities, accessibility, and control



Choosing the right one depends on your  
use case, security needs, and infrastructure

# Closed-source LLMs (Cloud-based APIs)

## 📌 Examples

- OpenAI (ChatGPT), Anthropic (Claude), Google (Gemini), Microsoft (Copilot)

## ✓ Advantages:

- Powerful and well-trained (best models available)
- Easy to use via APIs
- Regularly updated & improved

## ✗ Challenges:

- Black box (you don't control how they work)
- Expensive (API calls can add up quickly)
- Data privacy concerns (sending requests to external servers)

## 💡 When to use?

- If you need the most advanced models and don't mind API costs or external dependencies



# Open-source LLMs (Self- or cloud-hosted)

## 📌 Examples

- Meta's Llama 3, Mistral, Google's Gemma, Alibaba's Qwen

## ✓ Advantages:

- Greater control (you know exactly how the model works)
- Can be fine-tuned for specific needs
- No external API costs

## ✗ Challenges:

- Requires more setup (you have to run the model yourself)
- May not be as powerful as the latest closed models
- Needs infrastructure (e.g., GPUs for hosting)

## 💡 When to use?

- If you need control over the model & lower costs but are okay with slightly weaker performance



# Local models (on your machine or server)

## 📌 Examples

- Ollama, GGUF-based models (e.g., Llama, Mistral, Mixtral)

## ✓ Advantages:

- Works offline (great for security-sensitive applications)
- No API costs (completely free to use once set up)
- Low latency (responses are instant if hardware is good)

## ✗ Challenges:

- Limited by your hardware (needs a strong CPU/GPU)
- Not always as capable as cloud-hosted models
- Setup complexity (installing and optimizing models)

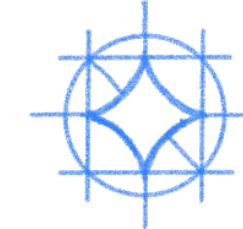
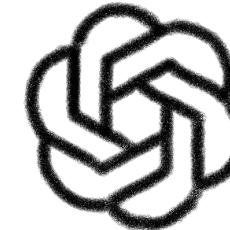
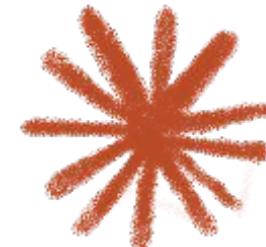
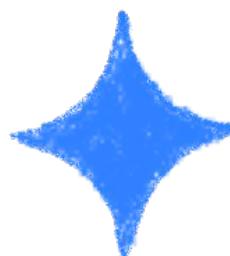
## 💡 When to use?

- If you need privacy and control, and you have the hardware to run an LLM efficiently



# Choosing the Right Model for your Apps

- Cloud APIs
  - Great for rapid development, but costly and not always secure
- Self-hosted open models
  - Best balance for long-term control and scalability
- Local models
  - Best for privacy-sensitive applications



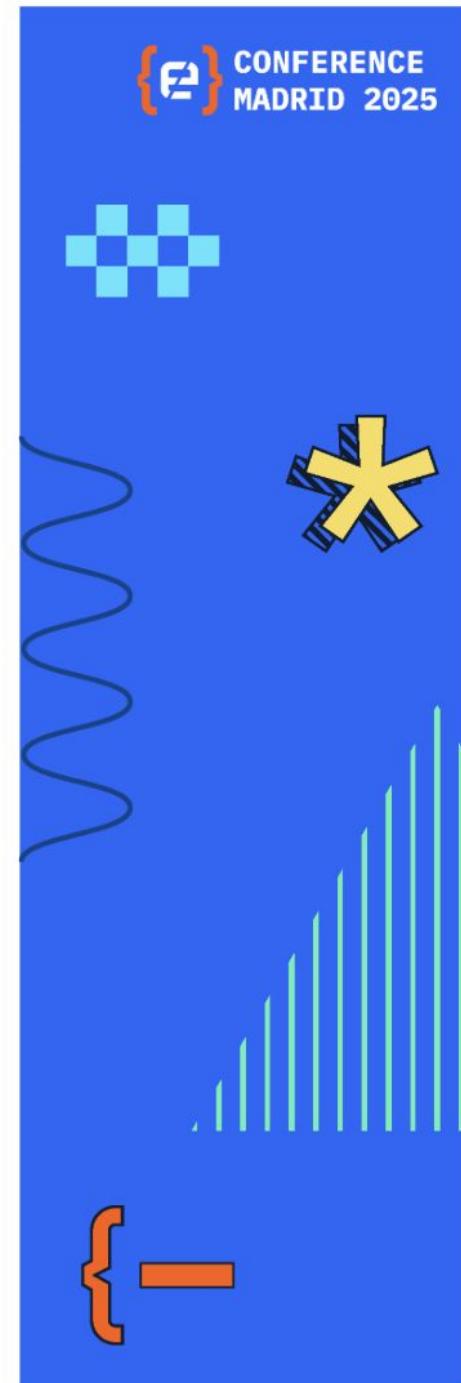
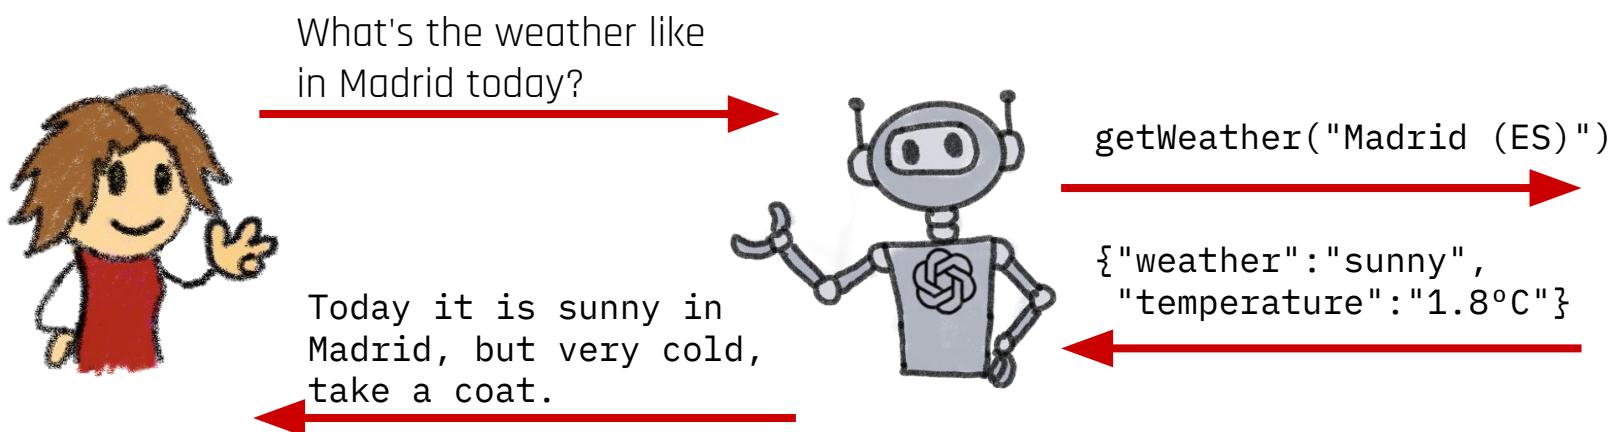
clever cloud



@LostInBrittany

# LLM evolution

From simple chat to tool-enhanced agent!

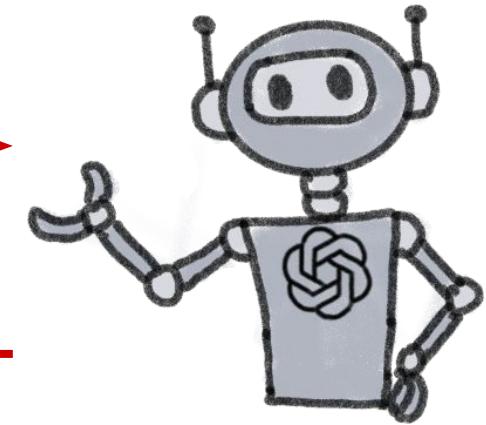


# LLM are only language models



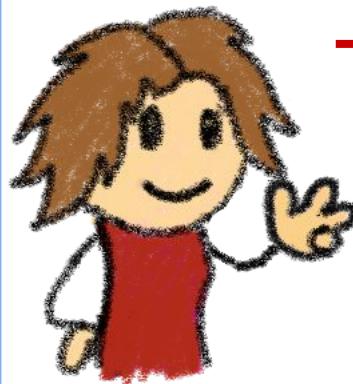
What's the weather like in Madrid today?

I'm unable to provide real-time information or current weather updates.

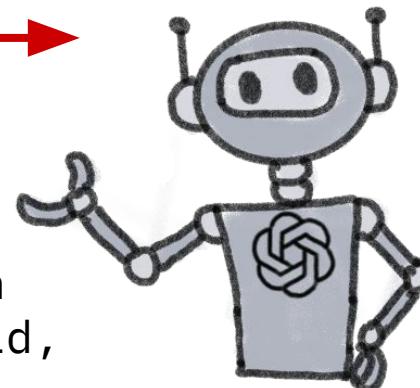


They have no built-in way to use external tools or real-time data

# Tools and plugins were added



What's the weather like  
in Madrid today?



Today it is sunny in  
Madrid, but very cold,  
take a coat.

getWeather("Madrid (ES)")

{ "weather": "sunny",  
"temperature": "1.8°C" }

Weather API



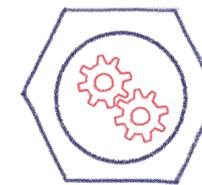
LLM recognizes it needs an external function  
and calls it, integrating the result into  
a natural-language response.

# LLM don't call directly those tools



What's the weather like in Madrid today?

AI Application

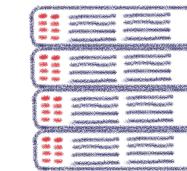


What's the weather like in Madrid today?  
If needed, you have an available weather  
tool: `getWeather(city)`



LLM

API



Call `getWeather("Madrid")`

`getWeather("Madrid")`

`{"weather": "sunny", "temperature": "1.8°C"}`

Result of the tool calling:  
`{"weather": "sunny", "temperature": "1.8°C"}`

Today it is sunny in Madrid, but very  
cold, take a coat.



clever cloud



@LostInBrittany

# How are those LLM Tools defined?



LyingWeatherTool.java

```
//DEPS dev.langchain4j:langchain4j:1.0.0-beta1

import dev.langchain4j.agent.tool.Tool;

public class LyingWeatherTool{
    @Tool("A tool to get the current weather in a city")
    public static String getWeather(String city) {
        return "The weather in " + city + " is sunny and hot.";
    }
}
```

# Why this matters?

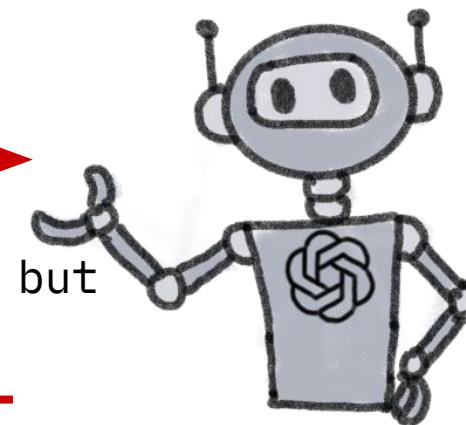
- Moves LLMs from static text generation
  - dynamic system components
- Increases accuracy & real-world usability
- Allows developers to control what the LLM can access



What's the weather like  
in Madrid today?



Today it is sunny in Madrid, but  
very cold, take a coat.



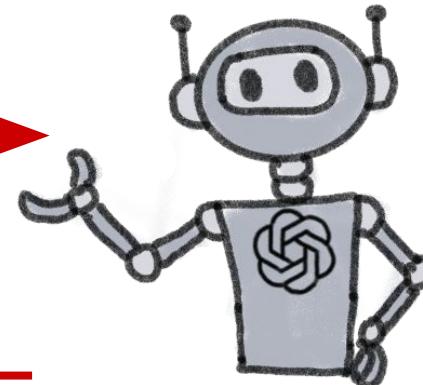
clever cloud

@LostInBrittany

# From LLM chats to LLM-powered agents



Can you summarize this  
YouTube video?



Of course, the video is a  
talk of Horacio about MCP...



YouTube



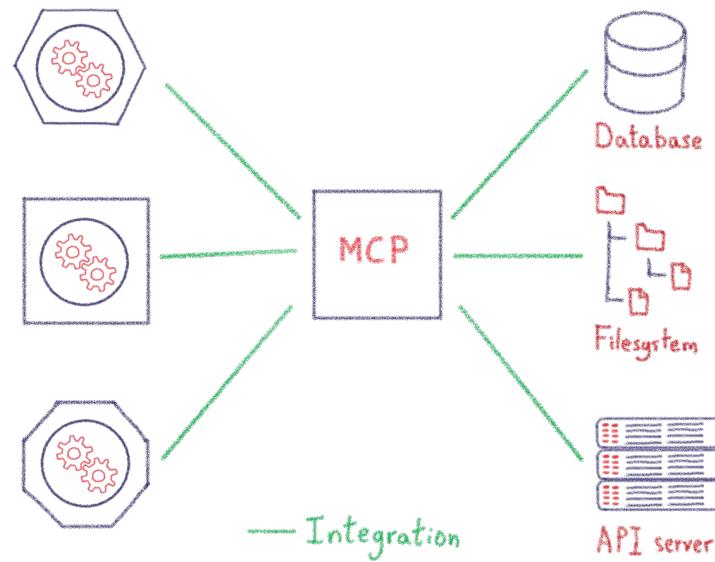
Whisper

\*This is a "fake" view, remember LLMs don't call tools directly  
But it's the view from the Point of View of the user

LLMs act like an agent that can plan actions:  
search the web, run some code, then answer

# Model Context Protocol (MCP): The missing link

MCP bridges LLMs with your applications,  
enabling controlled, real-world interactions



# Why Do We Need MCP?

Function calling is powerful,  
why do I need another concept?

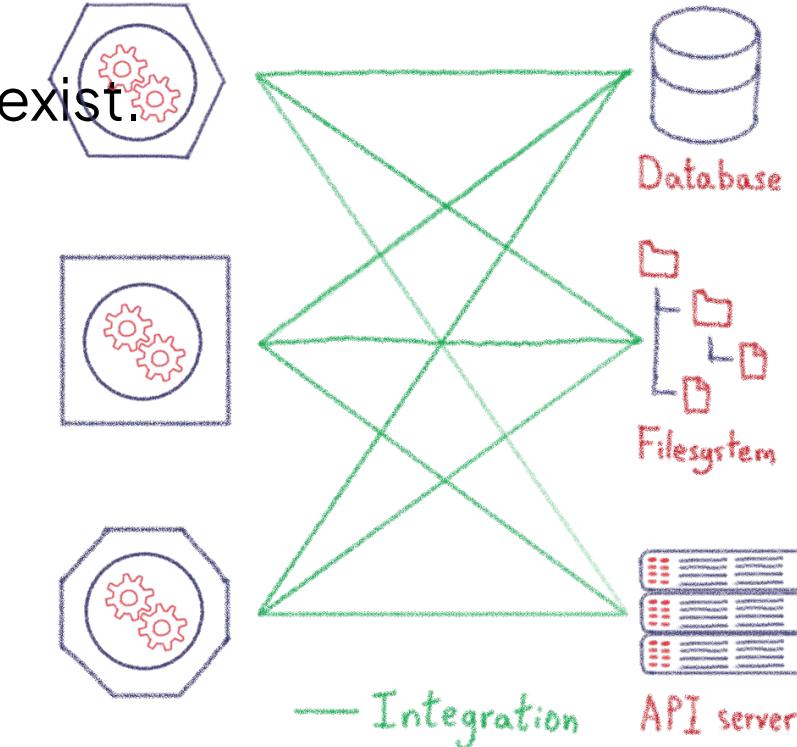


LLM function calling is useful,  
but it lacks structure

# Why Do We Need MCP?

## Problem

- LLMs **don't automatically know** what functions exist.
- **No standard way** to expose an application's capabilities.
- **Hard to control** security and execution flow.
- Expensive and fragile **integration spaghetti**

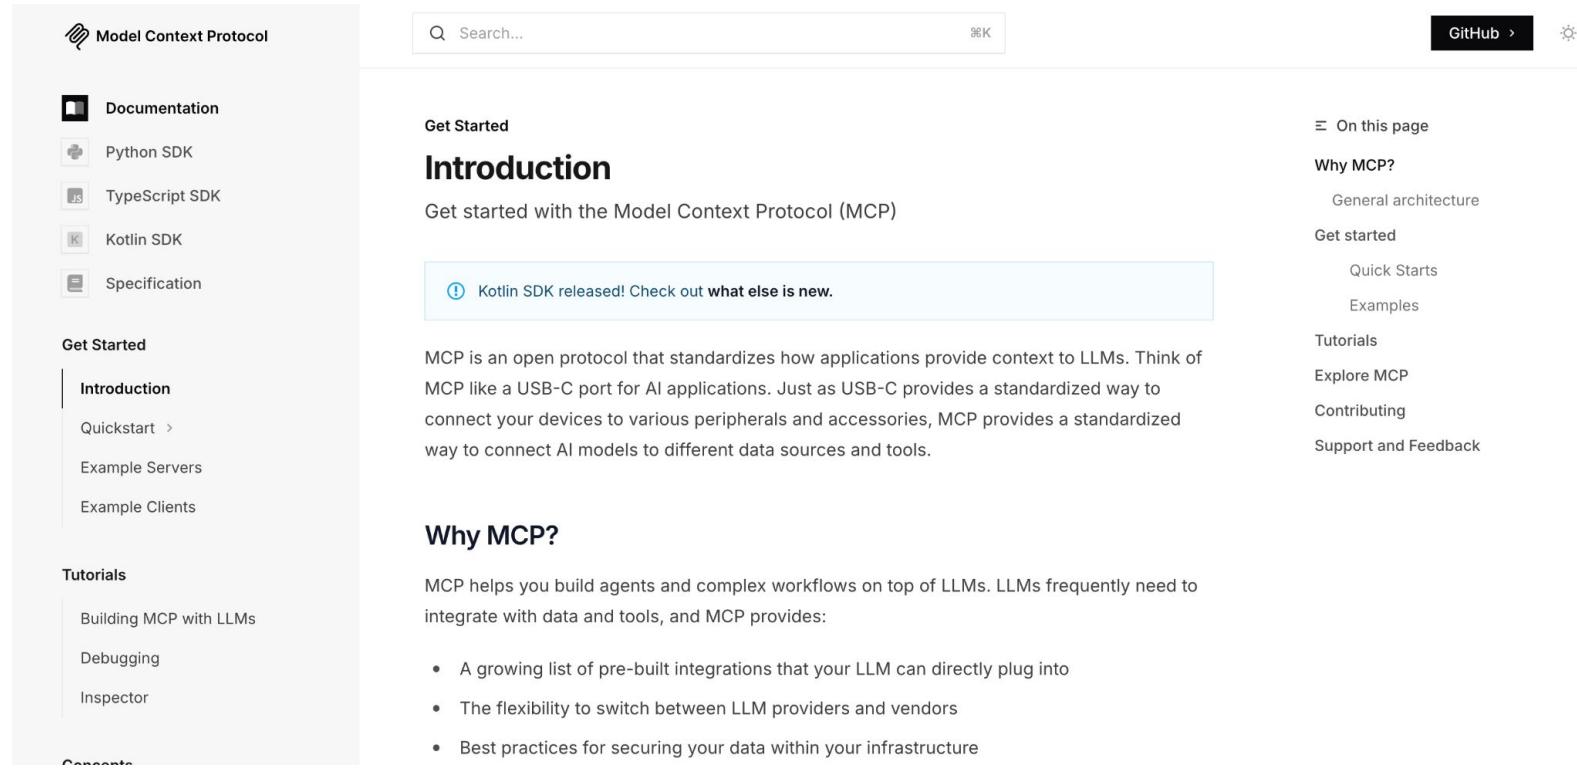


# Model Context Protocol



Anthropic, November 2024:  
*LLMs intelligence isn't the bottleneck,  
connectivity is*

# Model Context Protocol



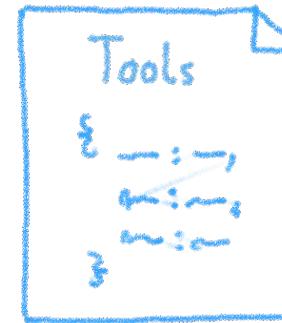
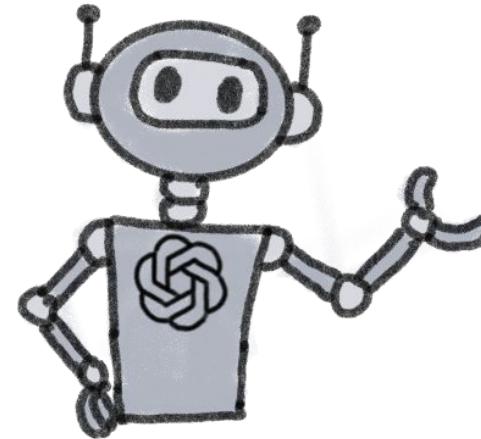
The screenshot shows the homepage of the Model Context Protocol (MCP) website. The header includes the MCP logo, a search bar, and navigation links for GitHub and the sun icon. The left sidebar has sections for Documentation (Python SDK, TypeScript SDK, Kotlin SDK, Specification), Get Started (Introduction, Quickstart, Example Servers, Example Clients), Tutorials (Building MCP with LLMs, Debugging, Inspector), and Concepts. The main content area features a "Get Started" section with an "Introduction" heading and a paragraph about MCP as an open protocol. A callout box says "Kotlin SDK released! Check out what else is new." Below this is a section titled "Why MCP?" with a list of benefits including pre-built integrations, flexibility, and best practices for security.

De facto standard for exposing system capabilities to LLMs

<https://modelcontextprotocol.io/>

# How MCP works

- Applications define an MCP manifest (structured JSON).
- The manifest describes available functions, input/output formats, and security policies.
- LLMs can discover and request function execution safely.

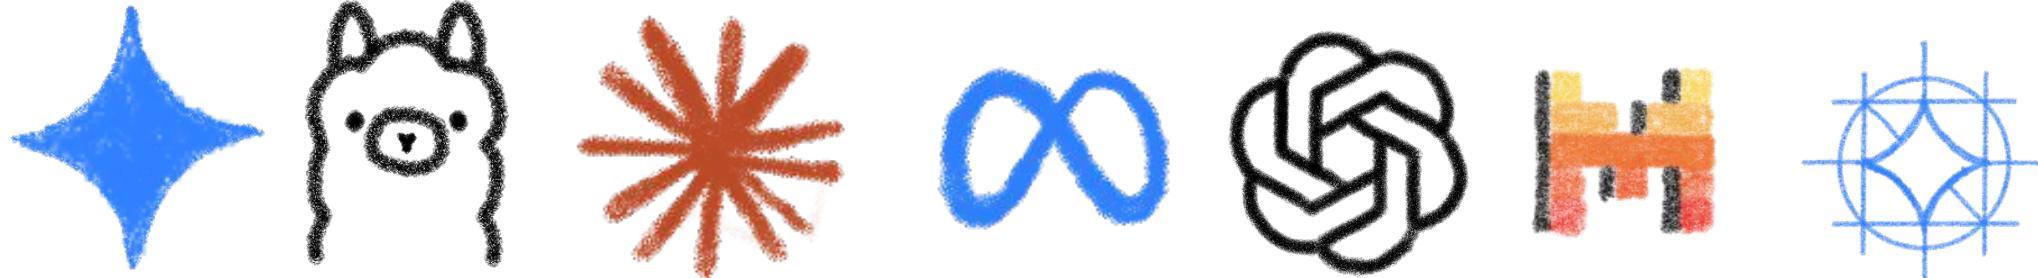


clever cloud



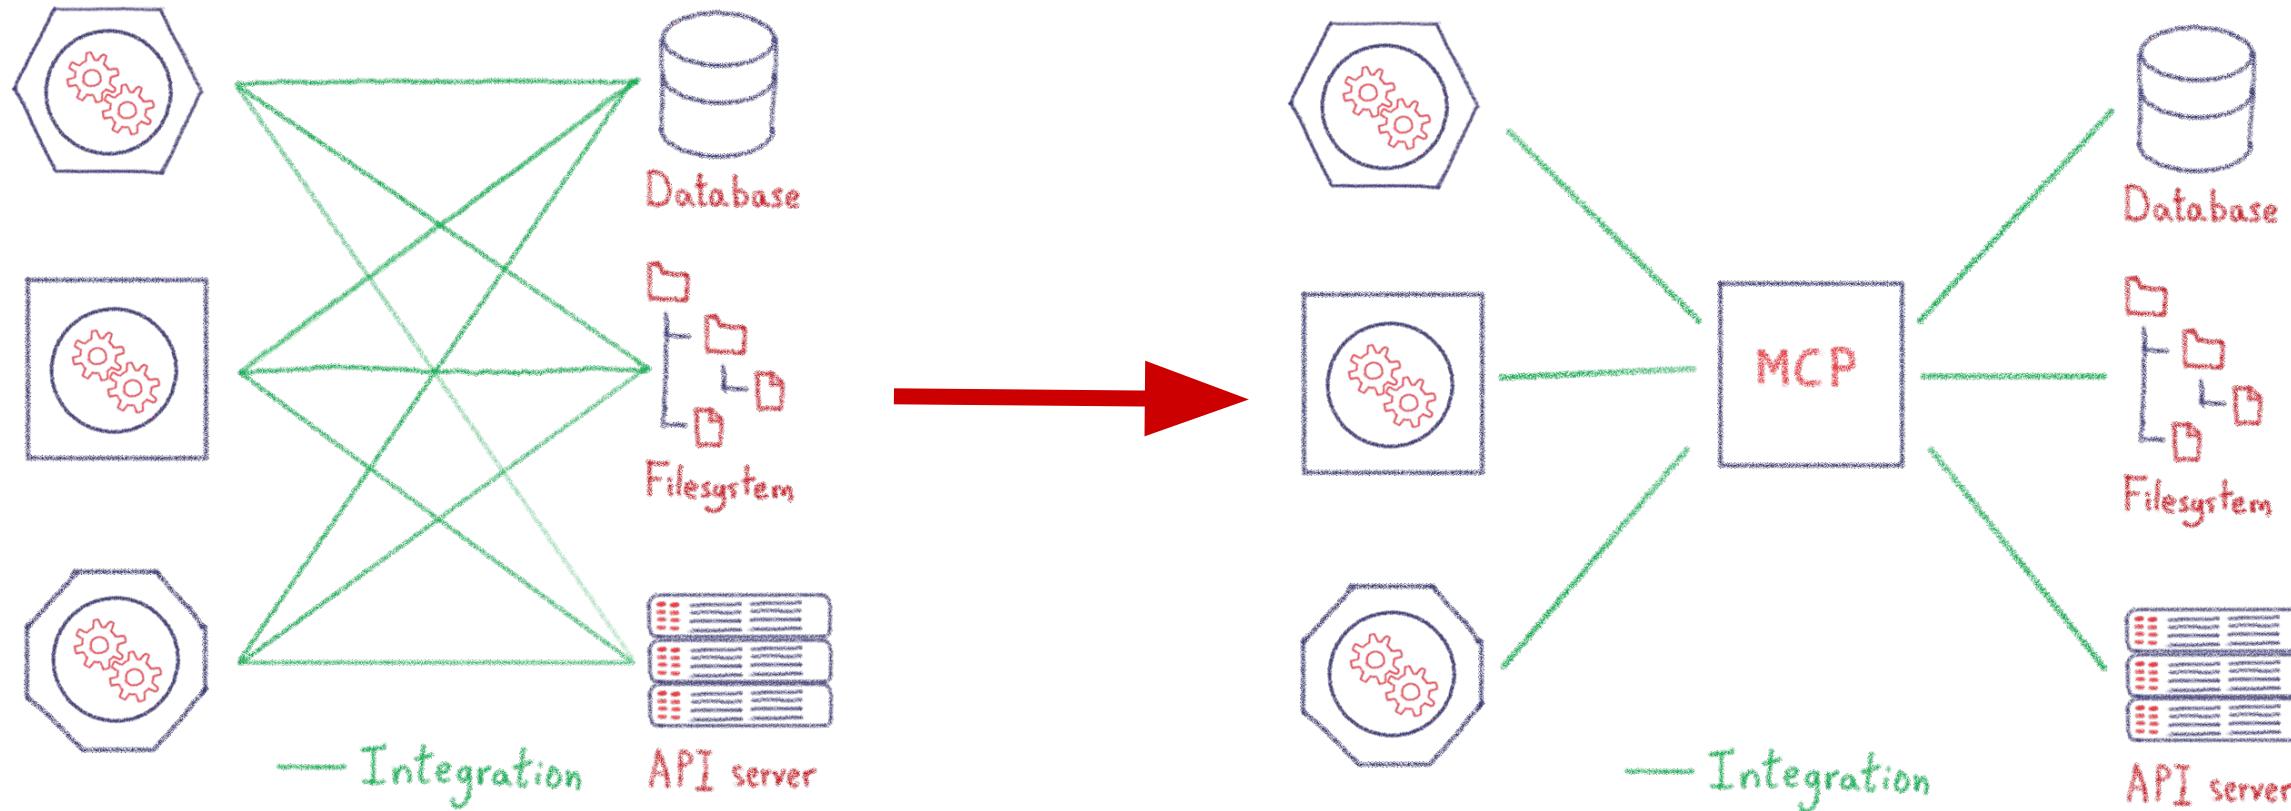
# MCP is provider-agnostic

Works with any LLM provider



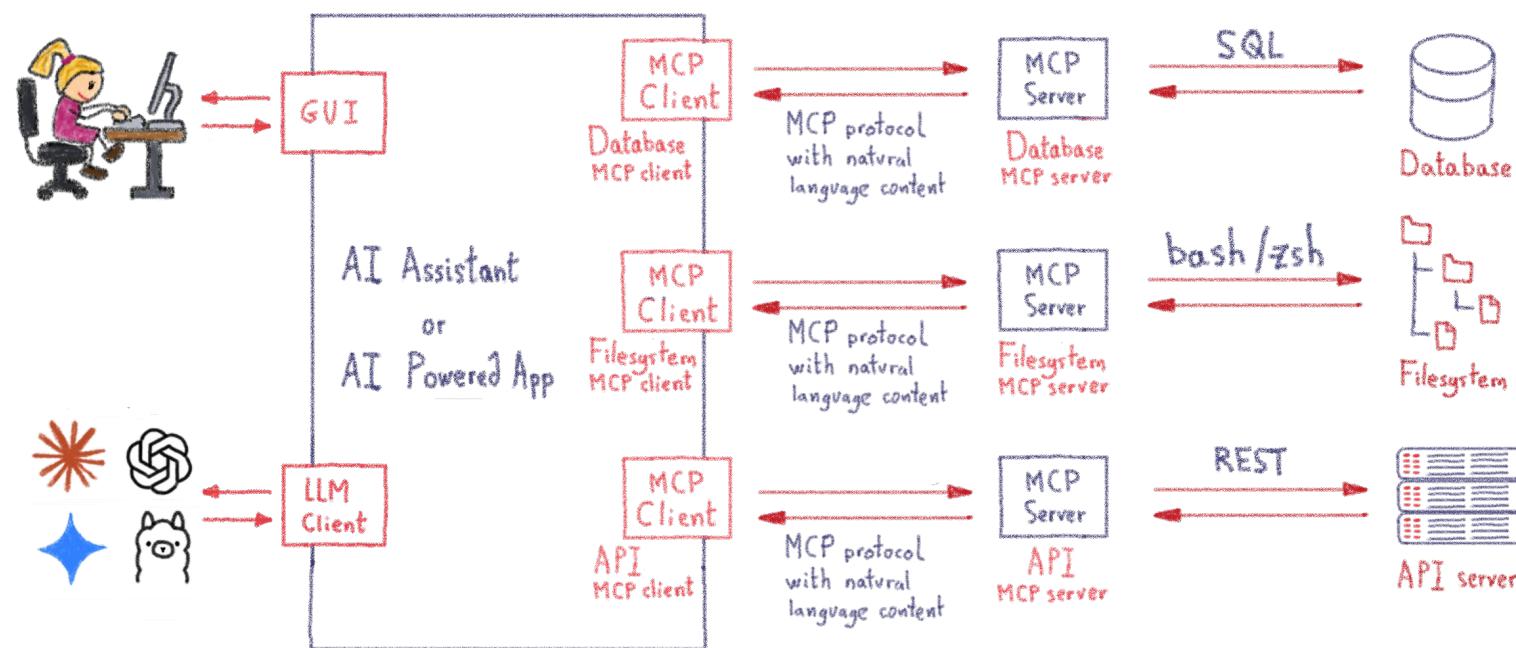
Ensures standardized function exposure  
across platforms

# MCP solves integration spaghetti

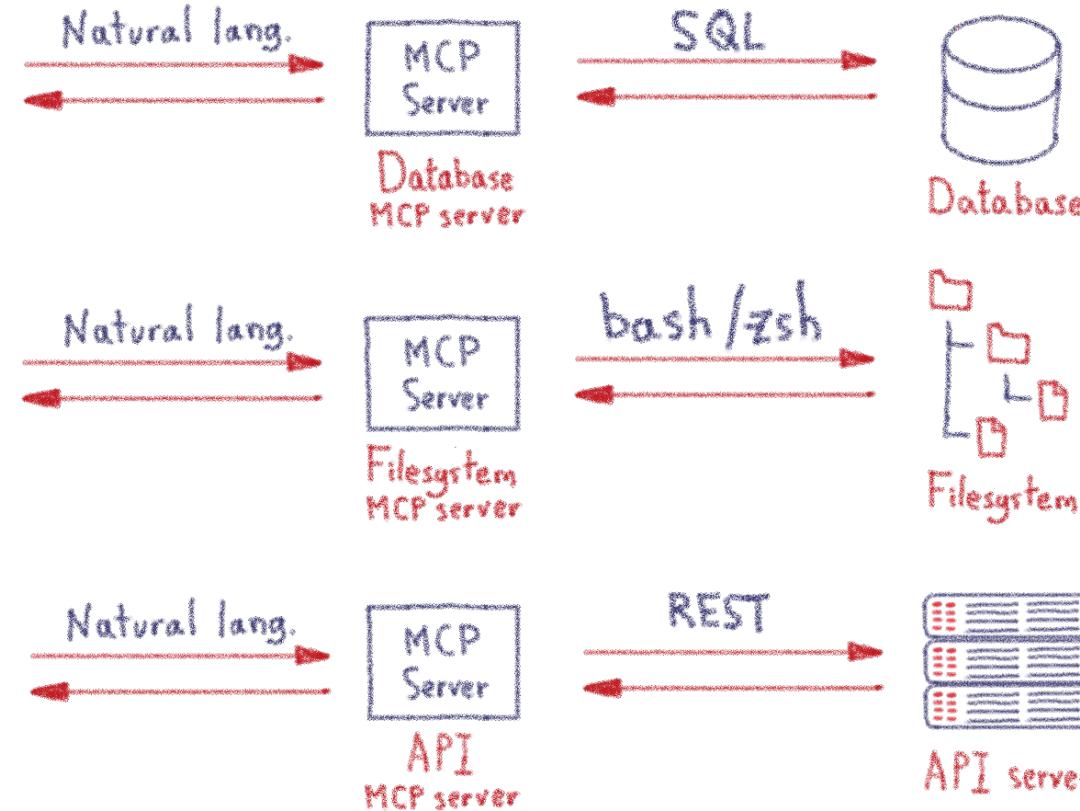


# The architecture of MCP

Clients, servers, protocol and transports  
Tools, resources and prompts

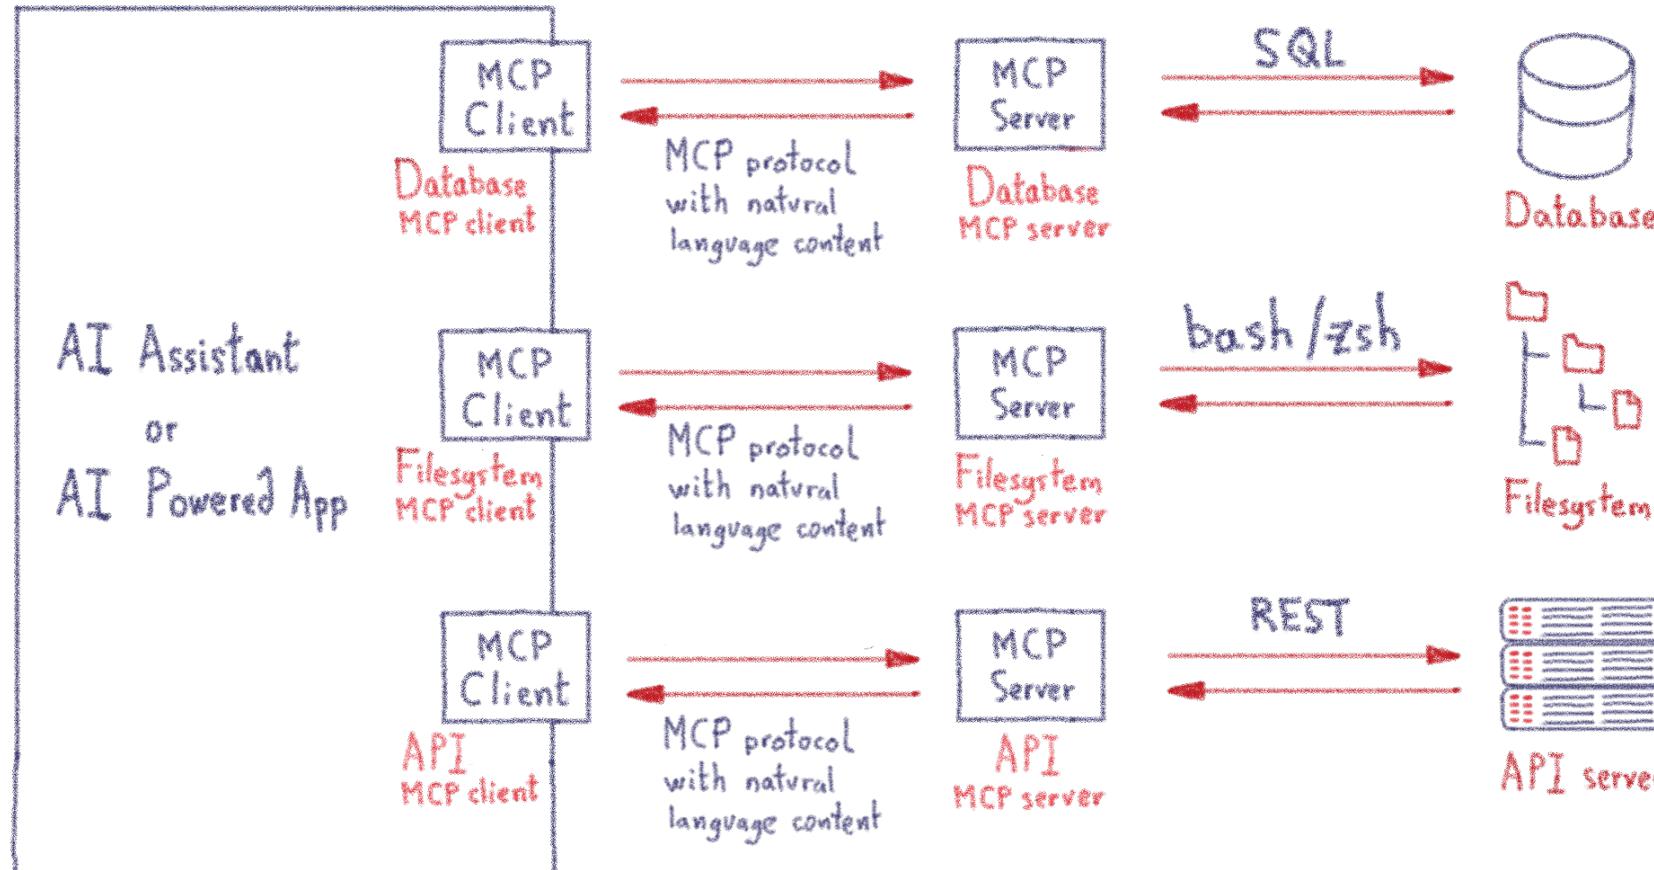


# MCP Servers: APIs in natural language



A new kind of API

# MCP Clients: on the AI assistant or app side



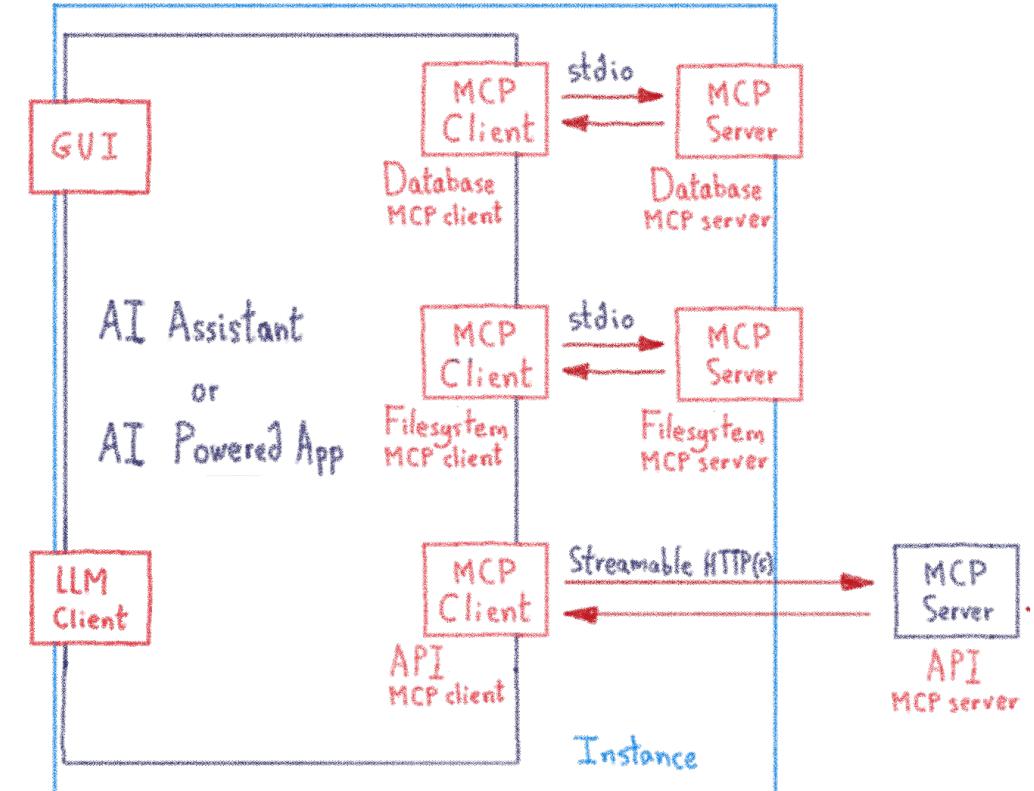
# MCP Protocol & Transports

## MCP Protocol

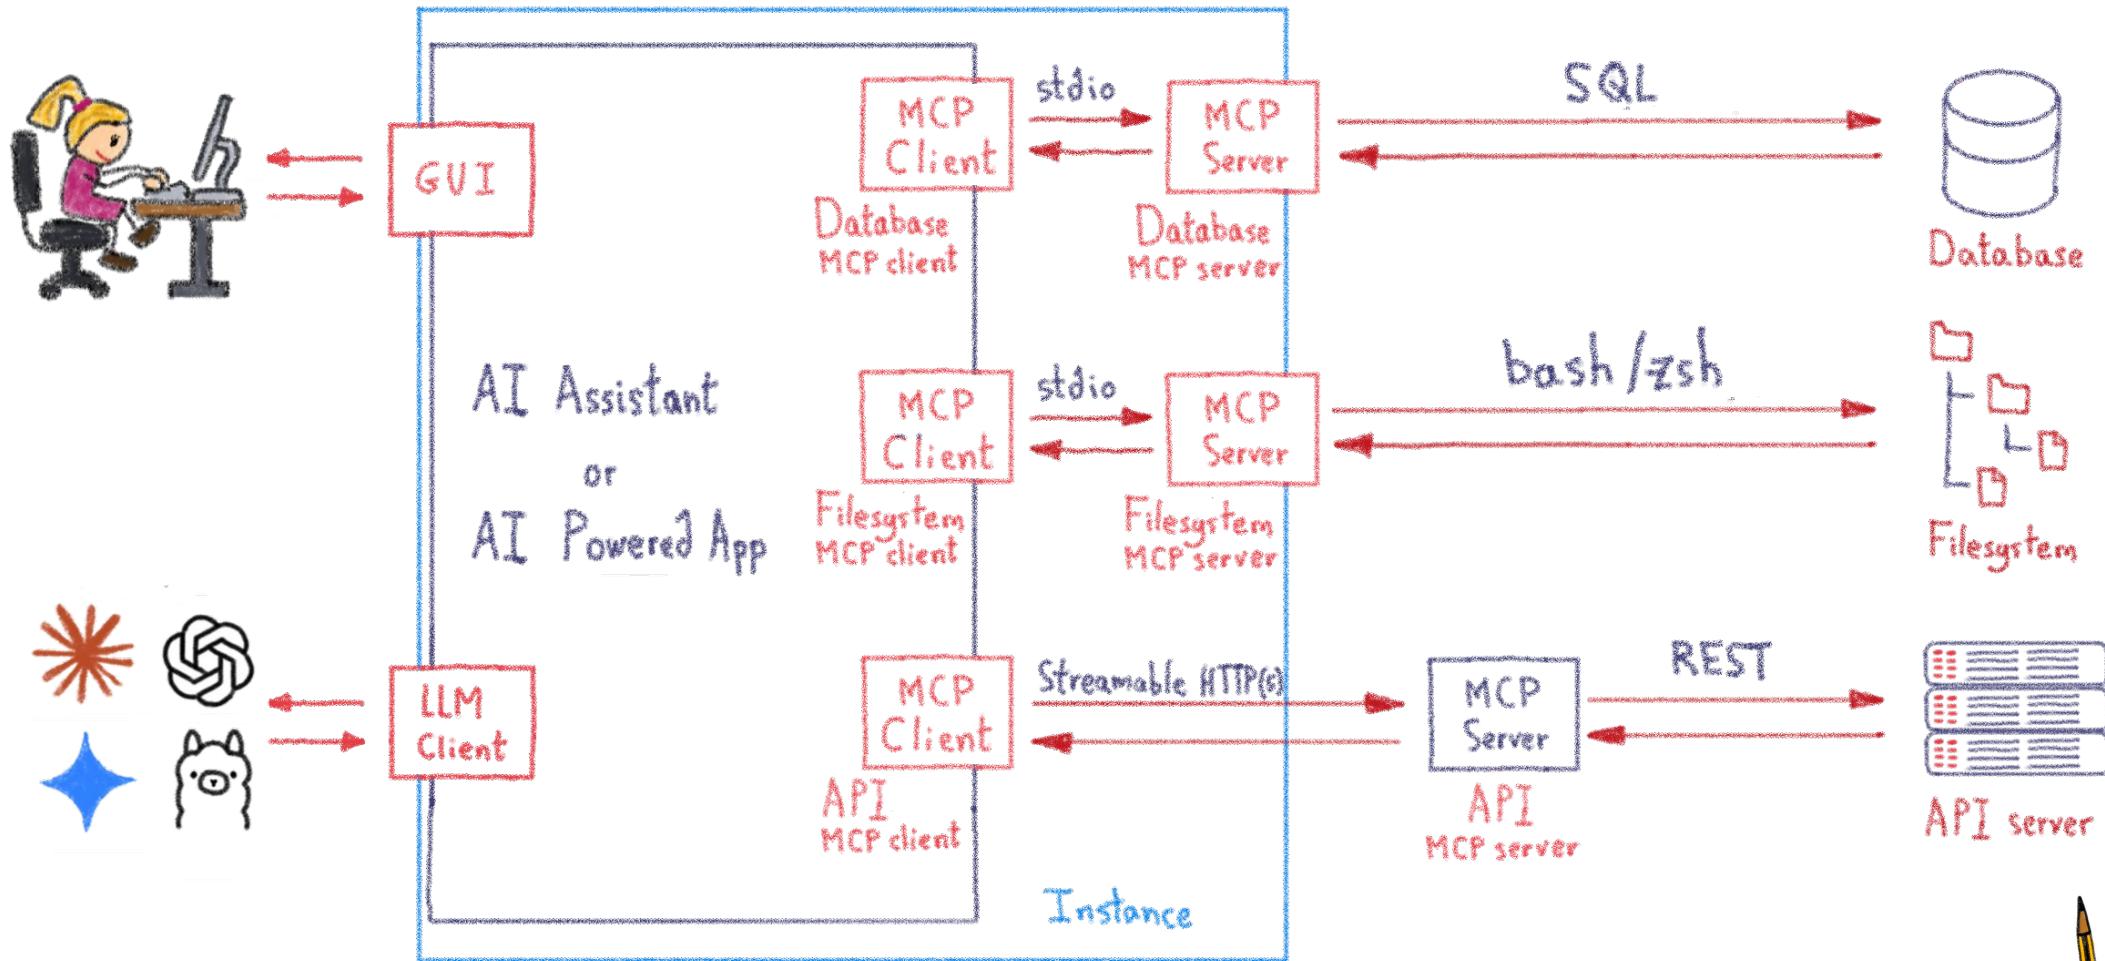
Follow the JSON-RPC 2.0 specification

## MCP Transports

- STUDIO (standard I/O)
  - Client and server in the same instance
- HTTP with SSE transport (deprecated)
- Streamable HTTP
  - Servers SHOULD implement proper authentication for all connections

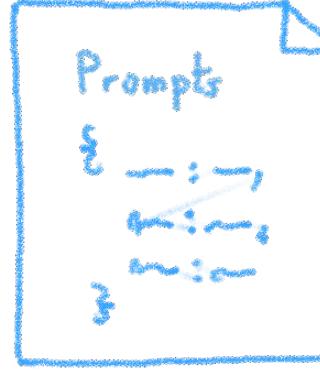
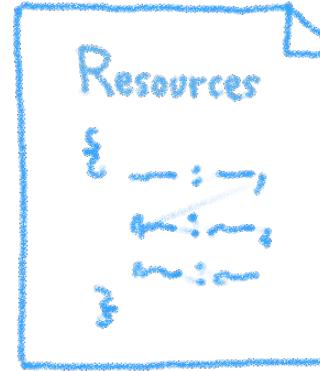
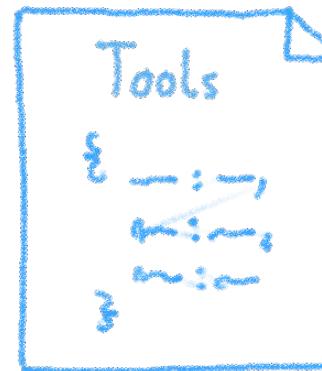


# Full MCP architecture

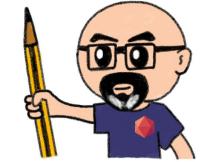
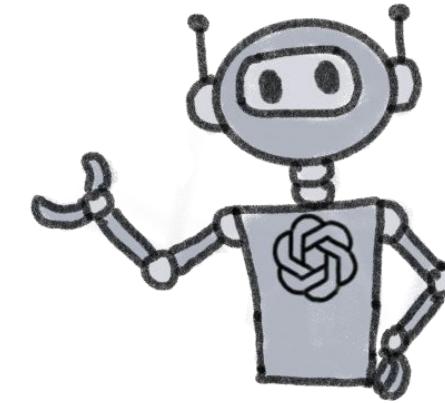


# Services: tools, resources & prompts

- Tools
  - Standardized way to expose functions that can be invoked by clients
- Resources
  - Standardized way to expose resources to clients
  - Each resource is uniquely identified by a URI
- Prompts
  - Standardized way to expose prompt templates to clients
  - Structured messages and instructions for interacting with LLMs



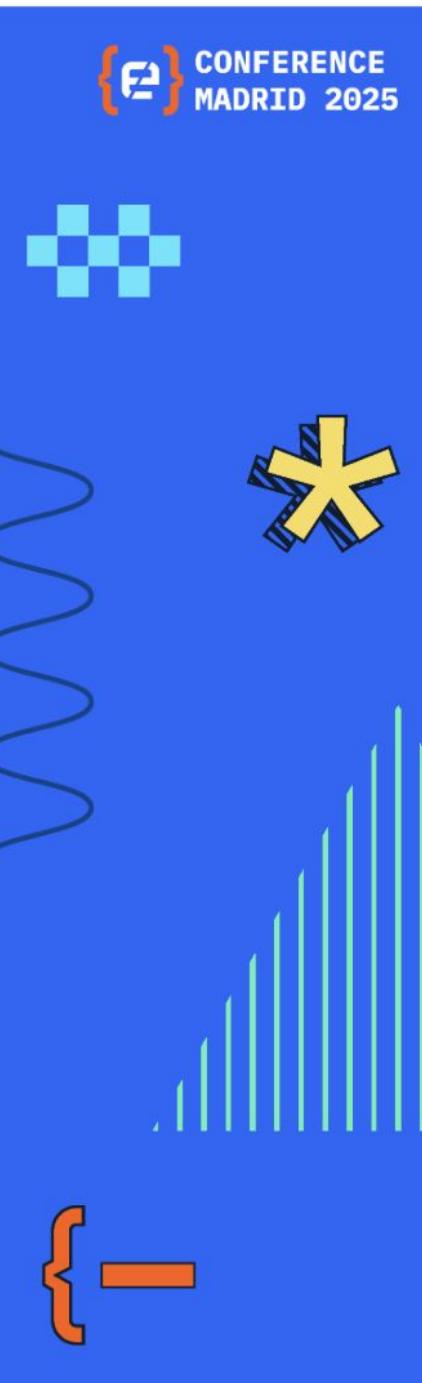
clever cloud



@LostInBrittany

# MCPs are APIs

And they should be architected in a similar way



# Let's use an example: RAGmonsters

README License

## RAGmonsters Dataset

### Overview

The RAGmonsters dataset is a collection of 30 fictional monsters created specifically for demonstrating and testing Retrieval-Augmented Generation (RAG) systems. Each monster is completely fictional and contains detailed information that would not be found in an LLM's training data, making it perfect for showcasing how RAG can enhance an LLM's knowledge with external information.

### Purpose

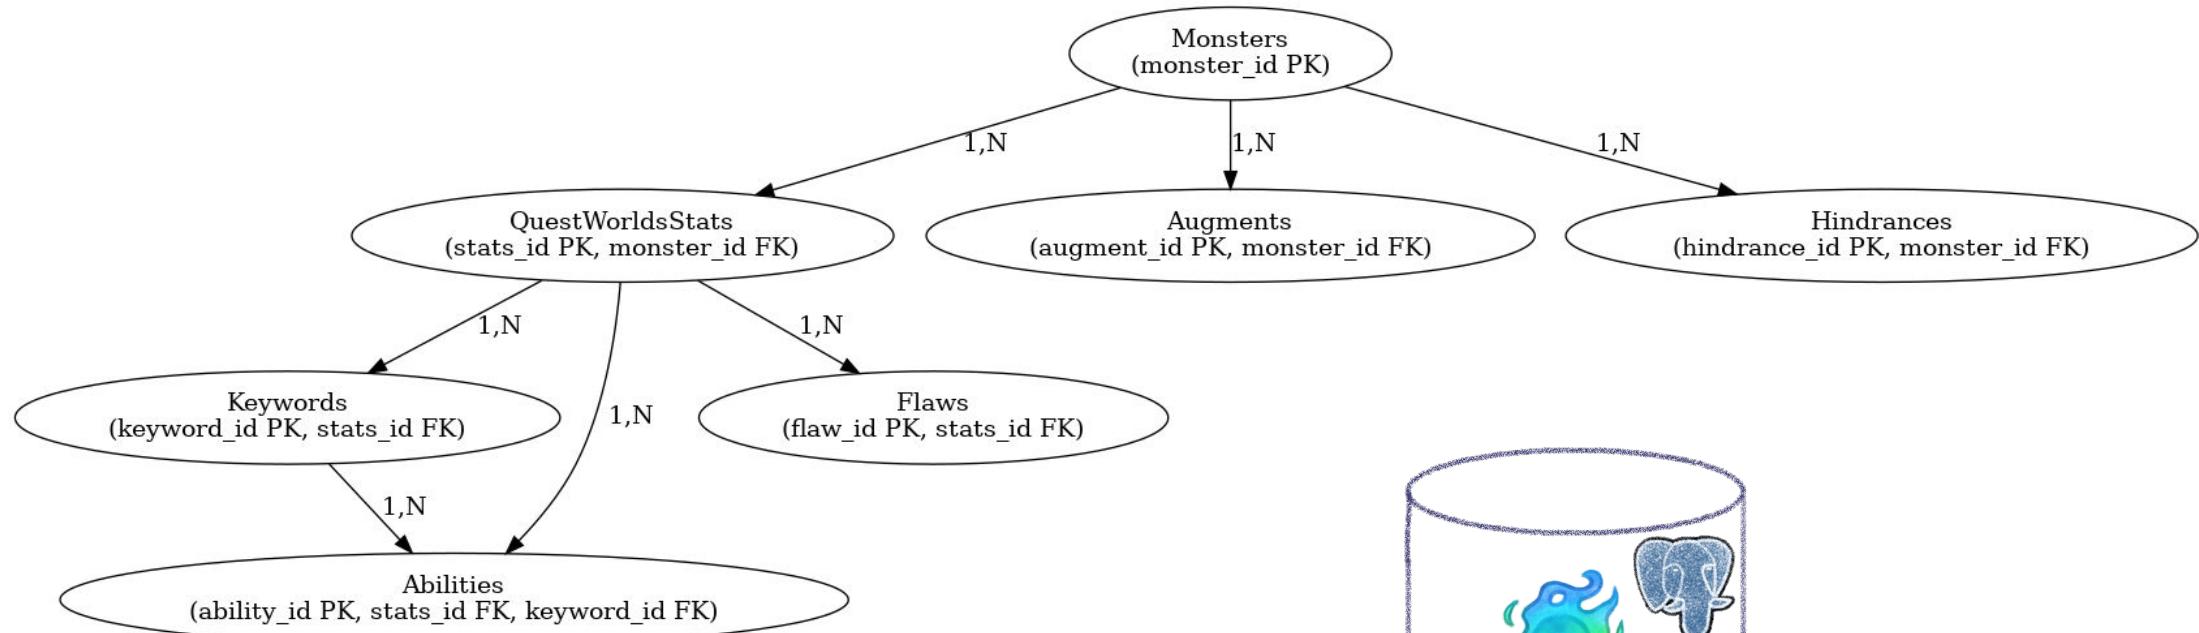
This dataset serves several educational purposes:

1. **Demonstrates RAG Value:** Shows how RAG can provide accurate answers about topics not in the LLM's training data
2. **Tests Retrieval Quality:** The varied attributes and relationships allow testing of different retrieval methods
3. **Supports Advanced Features:** Perfect for demonstrating filtering, re-ranking, and hybrid search techniques
4. **Provides Engaging Content:** Makes learning RAG concepts more fun and memorable



<https://github.com/LostInBrittany/RAGmonsters>

# RAGmonsters PostgreSQL Database

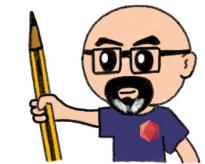
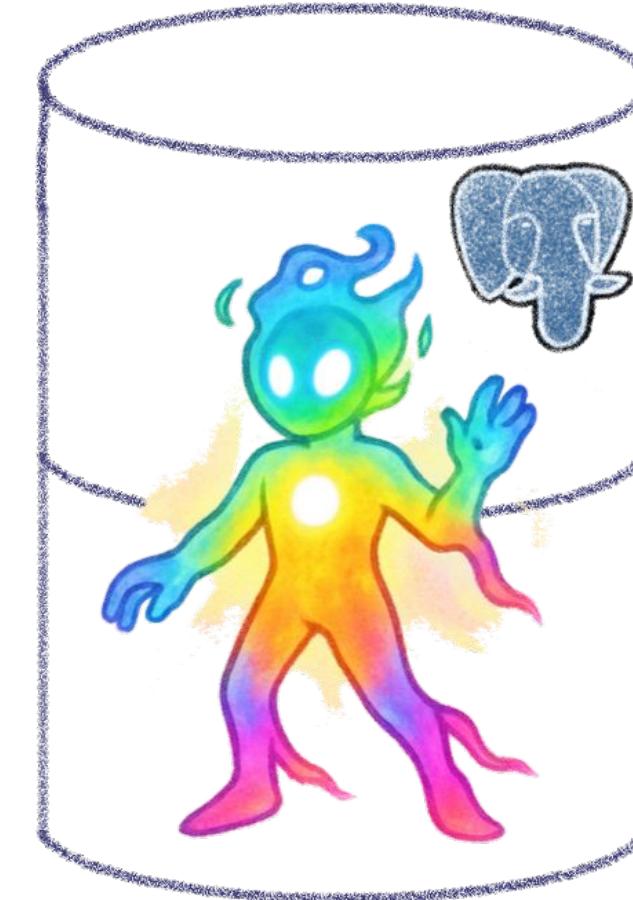


# We want to allow LLM request it

Two options:

- A generic PostgreSQL MCP server
- A custom-made MCP server tailored for RAGmonsters

Which one to choose?



@LostInBrittany

# Generic PostgreSQL MCP server

## Using PostgreSQL MCP Server

- A Resource that give the table schema for tables:  
[/schema](#)
- A Tool that allows to do SQL queries: [query](#)

LLM can know what tables do we have and what is their structure, and it can request them

Implementation:

<https://github.com/CleverCloud/mcp-pg-example>

PostgreSQL MCP Server:

<https://github.com/modelcontextprotocol/servers/tree/main/src/postgres>



### PostgreSQL MCP Chat

Ask questions about your PostgreSQL database in natural language

Hello! I'm your PostgreSQL assistant. Ask me anything about your database, and I'll help you explore it.

list the tables, please

Here are the tables in the database: 1. pg\_buffercache 2. pg\_stat\_statements\_info 3. spatial\_ref\_sys 4. geometry\_columns 5. combat\_statistics 6. flaws 7. abilities 8. augments 9. hindrances 10. keywords 11. monsters 12. questworlds\_stats 13. pg\_stat\_statements 14. raster\_columns 15. raster\_overviews 16. geography\_columns

Ask a question about your database...

Send



@LostInBrittany

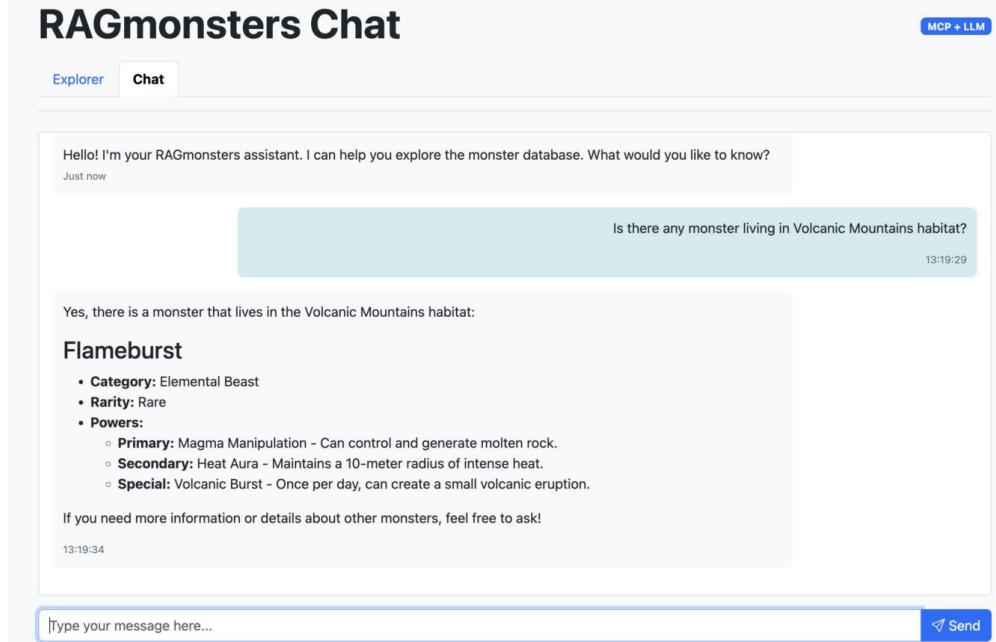
# Custom-made RAGmonsters MCP server

Coding a MCP server for it. It offers targeted tools:

- `getMonsterByName`: fetches detailed information about a monster.
- `listMonstersByType`: Lists monsters of a given type.
- Easy, intuitive interactions for LLMs.
- Optimized for specific use cases.
- Secure (no raw SQL).

Implementation:

<https://github.com/LostInBrittany/RAGmonsters-mcp-pg>



# How to choose?

Aspect	Generic MCP Server	Domain-Specific MCP Server
<b>Setup Speed</b>	Fast, minimal configuration	Slower, requires planning
<b>Efficiency</b>	Lower, LLM must explore schema	High, optimized for specific tasks
<b>Security</b>	Risk of SQL injection	Secure, predefined tools
<b>Flexibility</b>	Adapts to any schema	Needs updates with schema changes
<b>User Experience</b>	Complex, LLM must learn	Simple, guided interactions



# Conclusion

- Generic MCP servers: Quick to set up, flexible, but less efficient and more error-prone.
- Domain-specific MCP servers: Safer and faster for targeted tasks, but need more upfront design.
- Choose wisely: Use generic for exploration, domain-specific for production.

A bit like for REST APIs, isn't it?



# That's all, folks!

Thank you all!

